Machine Learning Engineer Nanodegree

Capstone Proposal

Simon Jackson March 20th, 2017

Proposal

Domain Background

The use of recommender systems is popular on e-commerce sites, which are interested in predicting the ratings/preferences of users for products[1] to personalise and improve the user experience. For example, Netflix wants to suggest movies you will like, Amazon ought to recommend books you're likely to enjoy, and so on.

Recommender systems often take the form of collaborative or content-based filtering. The former involves predicting what a user will like based on their similarity to other users[2]. The latter involves matching content to a user based on similarities among the content[3]. A more advanced approach that will be investigated in this project is a hybrid recommender system, which leverages and combines collaborative and content-based filtering.

In this project I will train a hybrid recommender system to predict how different users will rate (out of 5-stars) new movies. This system will be trained by combining information from the MovieLens 20M and IMDB 5000 Movie datasets, which are both available via the online machine learning challenge platform, Kaggle.

Problem Statement

Imagine you work at Netflix and have added new movies to the service. You'd like to recommend these to people who will like them. To determine whether a particular user might like one of these movies, it's impossible to see if other, similar users like the movie (because it's new and hasn't been rated). It's also challenging to see how the user rated other similar movies, because user ratings are relatively sparse. This project will attempt to solve this problem by investigating a hybrid recommender system for predicting the 5-star rating a person will give a new, unrated movie. A collaborative filtering component will link a given user to other similar users, and movie preferences can be pooled over the entire group of similar users. This can be used to estimate the preferences of many already-existing movies, which can be used by a content-filtering component to make the final prediction.

In summary, the goal is to create a recommender system that will predict users' ratings of new movies.

Datasets and Inputs

The data used in this project will come from two open-source projects:

- MovieLens 20M Dataset: Over 20 Million Movie Ratings and Tagging Activities Since 1995
- IMDB 5000 Movie Dataset: 5000+ movie data scraped from IMDB website

The MovieLens dataset is to be used as the key source for the collaborative filtering component of the model. It contains individual user ratings of movies on a 5-star scale (with 5 being the best and 1 being the lowest).

The IMDB dataset is to be used as the key source for the content-based filtering component of the model. It contains public information about 5000 movies and includes the following variables:

> "movie_title" "color" "num_critic_for_reviews" "movie_facebook_likes" "duration" "director_name" "director_facebook_likes" "actor_3_name" "actor_3_facebook_likes" "actor_2_name" "actor_2_facebook_likes" "actor_1_name" "actor_1_facebook_likes" "gross" "genres" "num_voted_users" "cast_total_facebook_likes" "facenumber_in_poster" "plot_keywords" "movie_imdb_link" "num_user_for_reviews" "language" "country" "content_rating" "budget" "title_year" "imdb_score" "aspect_ratio"

Combined, these two data sets can be used to train and test a hybrid recommender model for predicting the ratings that users will give "new" movies.

Solution Statement

The solution should:

1. Take a person who has rated movies in the MovieLens dataset
2. Take a movie that hasn't appeared in the MovieLens dataset, but has features about it available in the IMDB dataset.
3. Return a predicted 5-star rating.

For this particular project, this solution will need to combine collaborative and content-based filtering systems as a hybrid recommender system, as neither approach alone will be sufficient.

Benchmark Model

Two benchmark models for estimating the solution (5-star ratings):

- The mean rating of all users for all movies.
- The mean rating of the user for whom a prediction is being made.

Evaluation Metrics

Given a data set of known movie ratings, models investigating this problem can be evaluated by the accuracy of their predictions. For this project, these ratings will be treated as a continuous variable. Therefore, an appropriate metric for evaluating model performance will be the root mean square error (RMSE). This has been the metric used in similar problems such as the famous Netflix Prize).

The RMSE is calculated by squaring the error terms (residuals) for predictions on a given set of data points, calculating the means of these, and taking the square root. A value of 0 indicates that all predictions perfectly match the true values. The more positive the value, the worse the performance. Therefore, the expectation is that a hybrid recommender system will have a lower RMSE (closer to zero) than the benchmark approaches described above.

Project Design

My expected approach will attempt to do the following:

- Develop *collaborative filtering model*: For a particular user, use unsupervised methods to find similar users and how they would rate the movies being considered. This may inolve clustering users into different segments via methods such as K-means, or Gaussian Mixture models.
- Develop *content-based filtering model*: For any new movie, use unsupervised methods to find similar (rated) movies and link these to user ratings.
- Combine information in a supervised manner to make a prediction.

The tasks involved in achieving these steps include:

- Downloading and preprocessing the MovieLens 20M and IMDB 5000 Movie datasets.
- Exploring and visualizing the data.
- Building and evaluating benchmark models.
- Implementing an unsupervised algorithm for identifying movies that are similar to eachother.
- Implementing an unsupervised algorithm for identifying users whose ratings are similar to eachother.
- Implementing a supervised algorithm that combines information across these sources to make predictions of new movie ratings.
- Analyzing and evaluating the results.
- Considering improvements to the model.

The final data pipeline is expected to take IMDB information about a new movie and ratings of other movies for a given user, and produce a predicted 5-star rating.

---

1. Francesco Ricci and Lior Rokach and Bracha Shapira, Introduction to Recommender Systems Handbook, Recommender Systems Handbook, Springer, 2011, pp. 1-35
2. Prem Melville and Vikas Sindhwani, Recommender Systems, Encyclopedia of Machine Learning, 2010.
3. R. J. Mooney & L. Roy (1999). Content-based book recommendation using learning for text categorization. In Workshop Recom. Sys.: Algo. and Evaluation.